

✓ 1. Problem Definition & Understanding

Problem- Telecom and subscription-based businesses need to identify which customers are likely to cancel their service. Preventing churn is more cost-effective than acquiring new users.

Goal: Predict whether a customer will churn based on usage and interaction data.

✓ 2. Load Data & Cleaning

```
1 import pandas as pd
2
3 # Load dataset
4 df = pd.read_csv('/content/drive/MyDrive/churn_sample.csv')
5
6 # Preview
7 print(df.head())
8 print(df.info())
9 print(df['Churn'].value_counts())
10
```

	CustomerID	Contract	SupportCalls	MonthlyBill	PaymentMethod	BillingIssues	\
0	C001	Monthly	5	120.00	CreditCard	1	
1	C002	Annual	1	80.00	UPI	0	
2	C003	Monthly	5	74.12	CreditCard	1	
3	C004	Annual	1	46.44	UPI	0	
4	C005	Monthly	4	105.61	DebitCard	0	

	DataUsageGB	TenureMonths	AutoPay	Churn
0	60.00	6	0	Yes
1	95.00	24	1	No
2	105.81	36	1	Yes
3	74.44	21	0	No
4	105.19	16	1	Yes

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1002 entries, 0 to 1001

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	CustomerID	1002 non-null	object
1	Contract	1002 non-null	object
2	SupportCalls	1002 non-null	int64
3	MonthlyBill	1002 non-null	float64
4	PaymentMethod	1002 non-null	object
5	BillingIssues	1002 non-null	int64
6	DataUsageGB	1002 non-null	float64
7	TenureMonths	1002 non-null	int64
8	AutoPay	1002 non-null	int64
9	Churn	1002 non-null	object

dtypes: float64(2), int64(4), object(4)

memory usage: 78.4+ KB

None

Churn

No 528

Yes 474

Name: churn, dtype: int64

✓ 3. Data Preprocessing

```
1 from sklearn.preprocessing import LabelEncoder
2
3 df.drop('CustomerID', axis=1, inplace=True) #Drop ID column – not useful for prediction
4
5 label_cols = ['Contract', 'PaymentMethod', 'AutoPay', 'Churn']
6 le_dict = {}
7
8 for col in label_cols:
9     le = LabelEncoder()
10    df[col] = le.fit_transform(df[col])
11    le_dict[col] = le
12
```




```
13 # Check for missing values
14 print(df.isnull().sum())
```

```
Contract      0
SupportCalls  0
MonthlyBill    0
PaymentMethod  0
BillingIssues  0
DataUsageGB    0
TenureMonths   0
AutoPay        0
Churn          0
dtype: int64
```

✓ 4. Exploratory Data Analysis (EDA)

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4
5 # Set style
6 sns.set(style="whitegrid")
7
8 # Compute correlation matrix
9 corr = df.corr(numeric_only=True)
10
11 # Mask upper triangle
12 mask = np.triu(np.ones_like(corr, dtype=bool))
13
14 # Set up the matplotlib figure
15 plt.figure(figsize=(10,6))
16 sns.heatmap(corr, mask=mask, annot=True, fmt=".2f", cmap='Spectral', linewidths=0.5, square=True)
17 plt.title("Correlation Matrix (Lower Triangle Only)")
18 plt.show()
19
20
21 plt.figure(figsize=(8,5))
22 sns.boxplot(x='Churn', y='MonthlyBill', data=df, palette='Set3')
23 sns.swarmplot(x='Churn', y='MonthlyBill', data=df, color='black', alpha=0.6)
24 plt.title("Monthly Bill vs Churn (Box + Swarm Overlay)")
25 plt.xticks([0, 1], ['No Churn', 'Churn'])
26 plt.show()
27
28
29
30
```






/tmp/ipython-input-411382552.py:22: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `le`

`sns.boxplot(x='Churn', y='MonthlyBill', data=df, palette='Set3')`



Monthly Bill Distribution by Churn (KDE Plot)

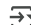
```

1 plt.figure(figsize=(8,5))
2 sns.kdeplot(data=df[df['Churn'] == 0], x='MonthlyBill', label='No Churn', shade=True)
3 sns.kdeplot(data=df[df['Churn'] == 1], x='MonthlyBill', label='Churn', shade=True)
4 plt.title("Monthly Bill Density by Churn Status")
5 plt.xlabel("Monthly Bill")
6 plt.legend()
7 plt.show()
8
9 avg_churn_bill = df[df['Churn'] == 1]['MonthlyBill'].mean()
10 avg_no_churn_bill = df[df['Churn'] == 0]['MonthlyBill'].mean()
11

```



```
12 print(f" Avg Monthly Bill (Churned): ₹{avg_churn_bill:.2f}")
13 print(f" Avg Monthly Bill (Retained): ₹{avg_no_churn_bill:.2f}")
```


 /tmp/ipython-input-2194561617.py:2: FutureWarning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

```
sns.kdeplot(data=df[df['Churn'] == 0], x='MonthlyBill', label='No Churn', shade=True)
/tmp/ipython-input-2194561617.py:3: FutureWarning:
```

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

```
sns.kdeplot(data=df[df['Churn'] == 1], x='MonthlyBill', label='Churn', shade=True)
```



Avg Monthly Bill (Churned): ₹90.37
Avg Monthly Bill (Retained): ₹90.71

✓ 5. ML Model Building

```
1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LogisticRegression
3
4 X = df.drop('Churn', axis=1)
5 y = df['Churn']
6
7 # Train-test split
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
9
10 # Train model
11 model = LogisticRegression(max_iter=1000) # Simple baseline model for binary classification
12 model.fit(X_train, y_train)
```



LogisticRegression ⓘ ?
LogisticRegression(max_iter=1000)

✓ 6. Model Evaluation

```
1 from sklearn.metrics import classification_report, confusion_matrix, f1_score
2 # Predictions
3 y_pred = model.predict(X_test)
4 # Evaluation
5 print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
6 print("\nClassification Report:\n", classification_report(y_test, y_pred))
7 print("F1 Score:", f1_score(y_test, y_pred))
```



```
Confusion Matrix:
[[89 69]
 [69 74]]
```

```
Classification Report:
              precision    recall  f1-score   support


     0       0.56       0.56       0.56       158
     1       0.52       0.52       0.52       143

 accuracy          0.54
 macro avg          0.54
 weighted avg       0.54
```

F1 Score: 0.5174825174825175

7. Feature Importance

```
1 coeffs = pd.Series(model.coef_[0], index=X.columns)
2 coeffs.sort_values().plot(kind='barh', figsize=(10,6), title='Feature Influence on Churn')
3 plt.xlabel("Coefficient Value")
4 plt.show()
```



8. Predict New Customer

```
1 new_customer = pd.DataFrame({
2     'Contract': [le_dict['Contract'].transform(['Monthly'])[0]],
3     'SupportCalls': [4],
4     'MonthlyBill': [110.0],
5     'PaymentMethod': [le_dict['PaymentMethod'].transform(['CreditCard'])[0]],
6     'BillingIssues': [0],
7     'DataUsageGB': [85.0],
8     'TenureMonths': [12],
9     'AutoPay': [le_dict['AutoPay'].transform([1])[0]]
10 })
11 pred = model.predict(new_customer) ## Predict churn for new customer input
12 print("Predicted Churn:", le_dict['Churn'].inverse_transform(pred)[0])
```

Predicted Churn: Yes

